



# **Embarrassingly Parallel Computations**

**CPSC 424/524**

**Lecture #7**

**October 10, 2018**




# Parallel Techniques

- Embarrassingly Parallel Computations
- Partitioning and Divide-and-Conquer Strategies
- Pipelined Computations
- Load Balancing and Termination Detection



# Embarrassingly Parallel Computations (EPCs)

A computation that can be divided into a number of nearly independent parts (tasks), each of which can be executed by a separate entity (thread, cpu, process).



- No (or very little) coordination among tasks
- Tasks have lots of computation relative to the amount of coordination or communication (including startup/shutdown activities)
- Often implemented using “master-worker” or “manager-worker” model



# EPC Examples

- Low-level image processing (e.g. rendering)
- Visualization of the Mandelbrot set
- Genomics Algorithms (e.g., BLAST search)
- Monte Carlo Computations



## Example: Mandelbrot Set

Set of points  $c$  in the complex plane for which the sequence of points  $z_0 (=0)$ ,  $z_1$ ,  $z_2$ , ... computed by:

$$z_{k+1} = z_k^2 + c$$

is “quasi-stable” (will increase and decrease, but not exceed some limit). It can be shown that if  $z_{k+1}$  ever exceeds 2, then  $c$  is not in the set. Computing the set requires testing many points.

Algorithm used for testing membership in the set:

For a given  $c$ , tentatively mark  $c$  as in the set. Then iterate:

for ( $k = 0$ ,  $z_k = 0$ .;  $k < k_{max}$ ;  $k++$ ) {

    Compute the magnitude of  $z_{k+1} = z_k^2 + c$ ;

    If  $|z_{k+1}| > 2$ , then mark  $c$  as not in the set and break;

}




# Visualizing the Mandelbrot Set


Superimpose a mesh on a region of the complex plane. Map each point  $c$  in the mesh to a pixel in a rectangular image, carry out the iteration for  $c$ , and assign its pixel a color according to:

$c$  in the set:                      Black

$c$  not in the set:                Color corresponding to index  $k$  for which  
 $|z_{k+1}| > 2$



# Static (“Up Front”) Task Mapping



Problem:

Amount of work per point is unpredictable & highly variable, although there are “good” and “bad” regions where nearby points require similar amounts of work.

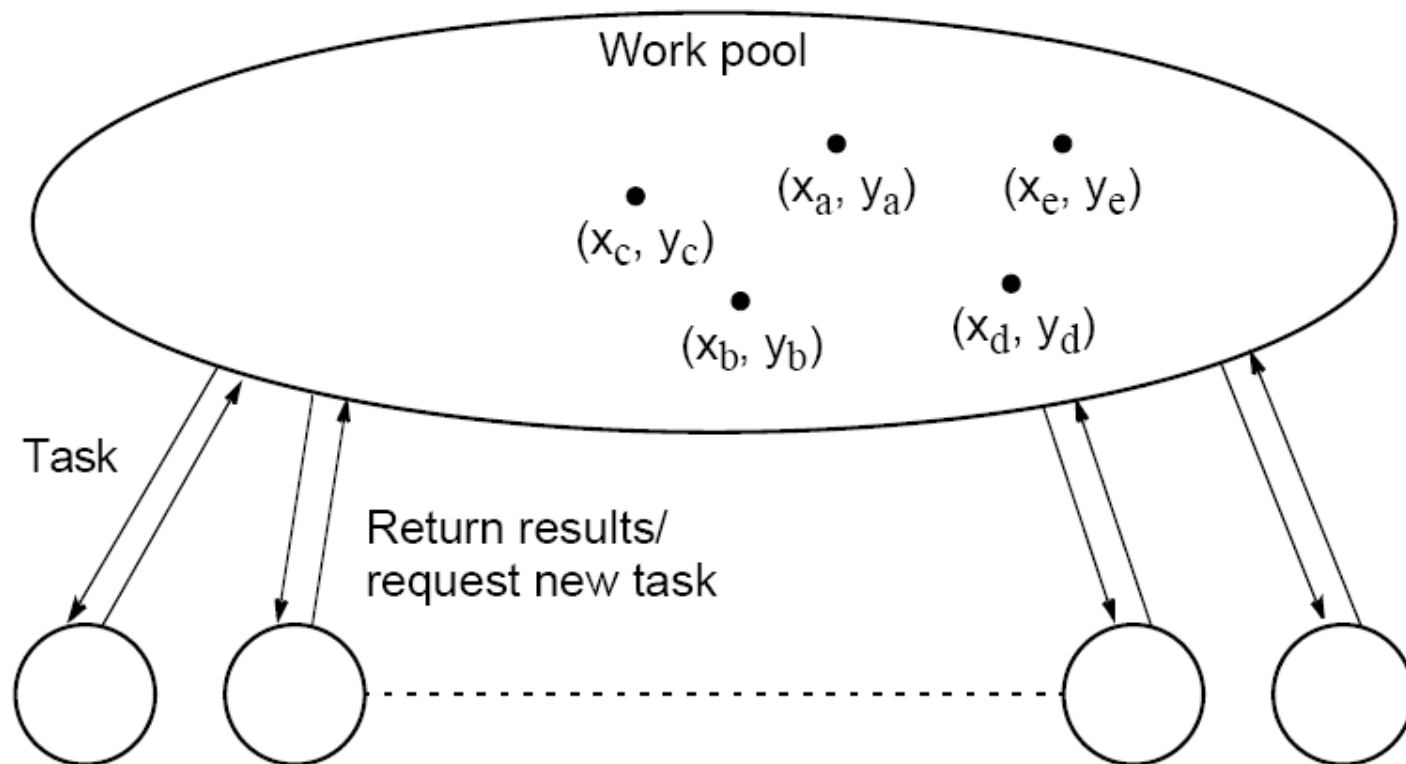
One solution:

Use static assignment with random mapping (by pixel)



# Dynamic (“On-the-Fly”) Task Assignment

Idea: Create many tasks ( $\gg$  number of workers), with each having a large **computation-to-coordination ratio**. Then give every worker a small number of tasks (often just one) and resupply them on demand.

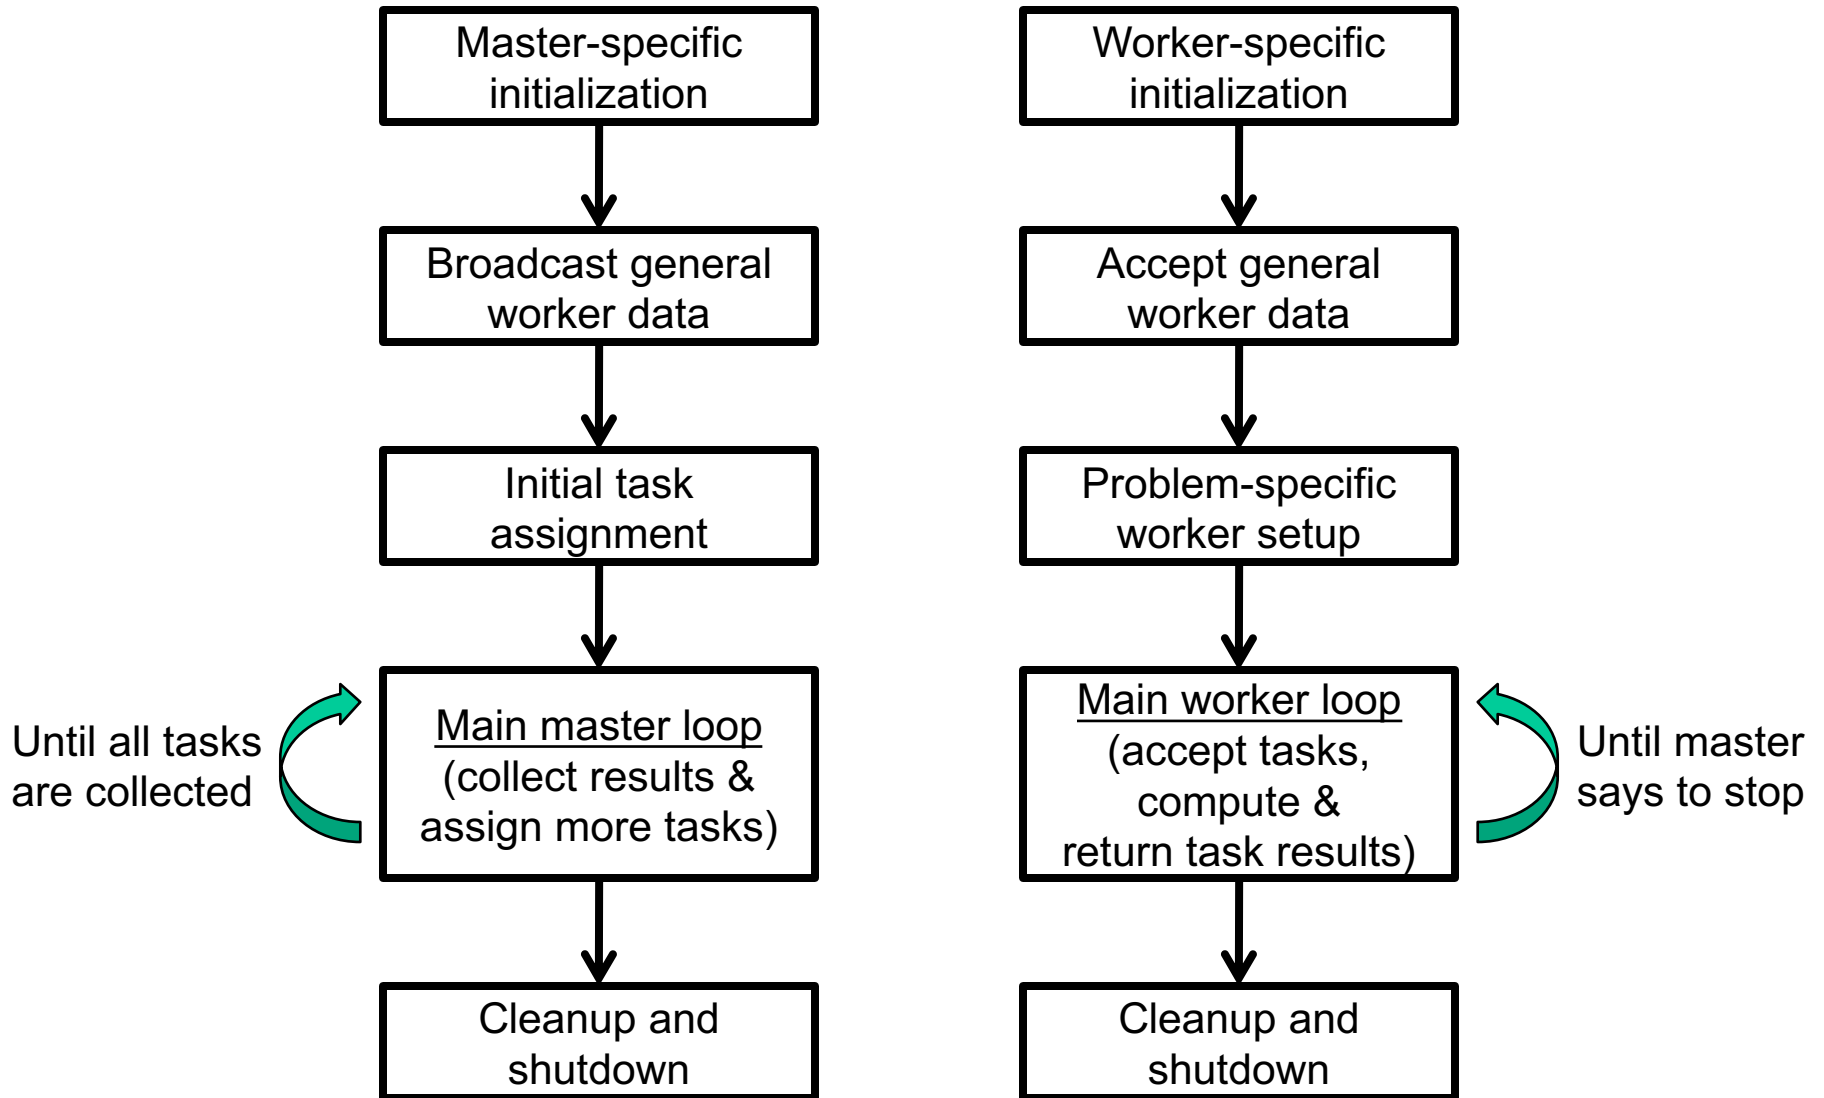


*(How well will this work for the Mandelbrot computation?)*





# Dynamic Task Pool Master/Worker

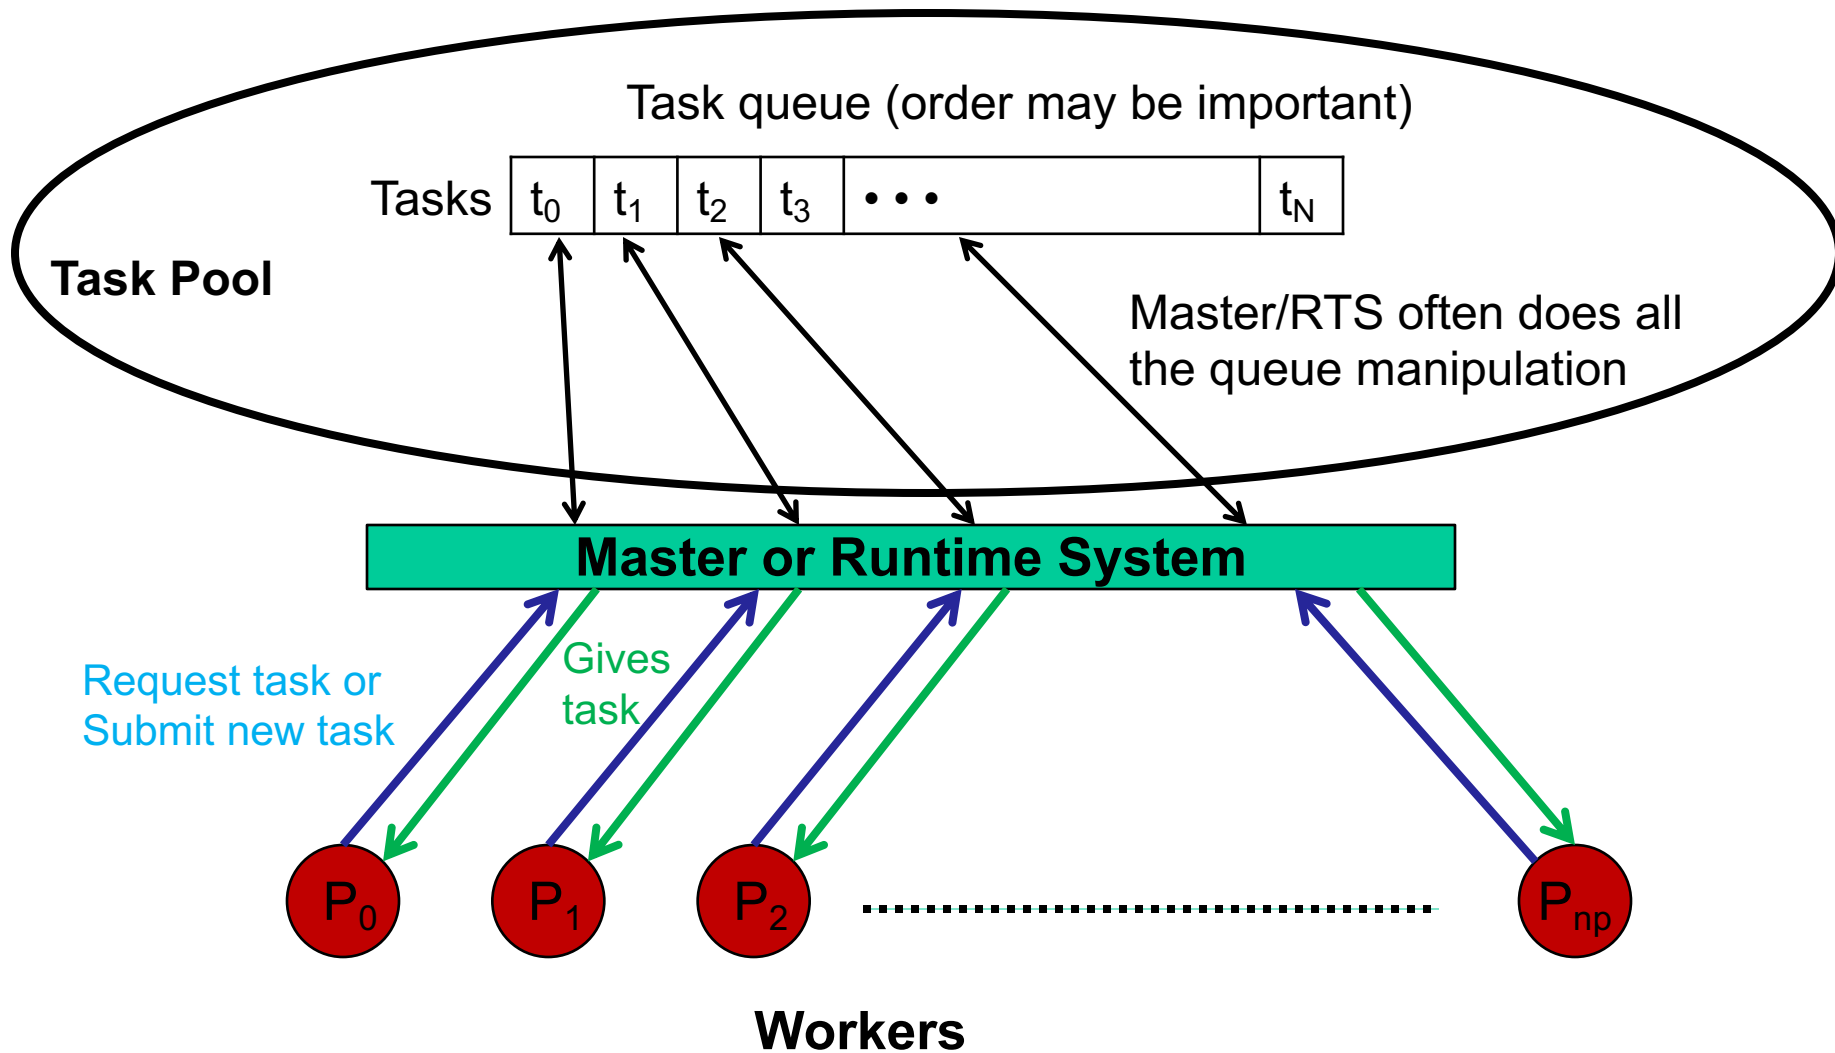


# Dynamic Load Balancing

- Centralized Task Pool Management
  - Master creates a task pool (“task bag”). Watermarking may be used.
  - Workers take or receive task assignments when idle. The master may do work or act purely as a manager to balance the load.
  - Workers simply process tasks and (in some applications) create new tasks that are added to the task pool.
  - Termination: Master (or runtime system) terminates when all tasks are done and workers are in states where no more tasks will be added.
- Decentralized
  - Master passes out all/initial tasks to workers
  - Workers cooperate to balance the load.
  - Master may behave as a worker in addition to its master role
  - Termination: Often more complex since no single process or thread really knows everything



# Centralized Task Pool Model



# Termination

- Task Queue must be empty **and**
- All tasks that will ever be created must be complete
  - For fixed task set: Master/RTS creates tasks and tracks completion of tasks and collection of results. That's what the Mandelbrot code does—either explicitly (e.g., via OMP loops) or implicitly (e.g., via task synchronization points).
  - For dynamic task set (when workers can submit new tasks):
    - All workers must be idle. (For example, they are all waiting for tasks when none is available.)
    - **Not sufficient** to terminate when task queue is empty if one or more workers might still be able to submit new tasks

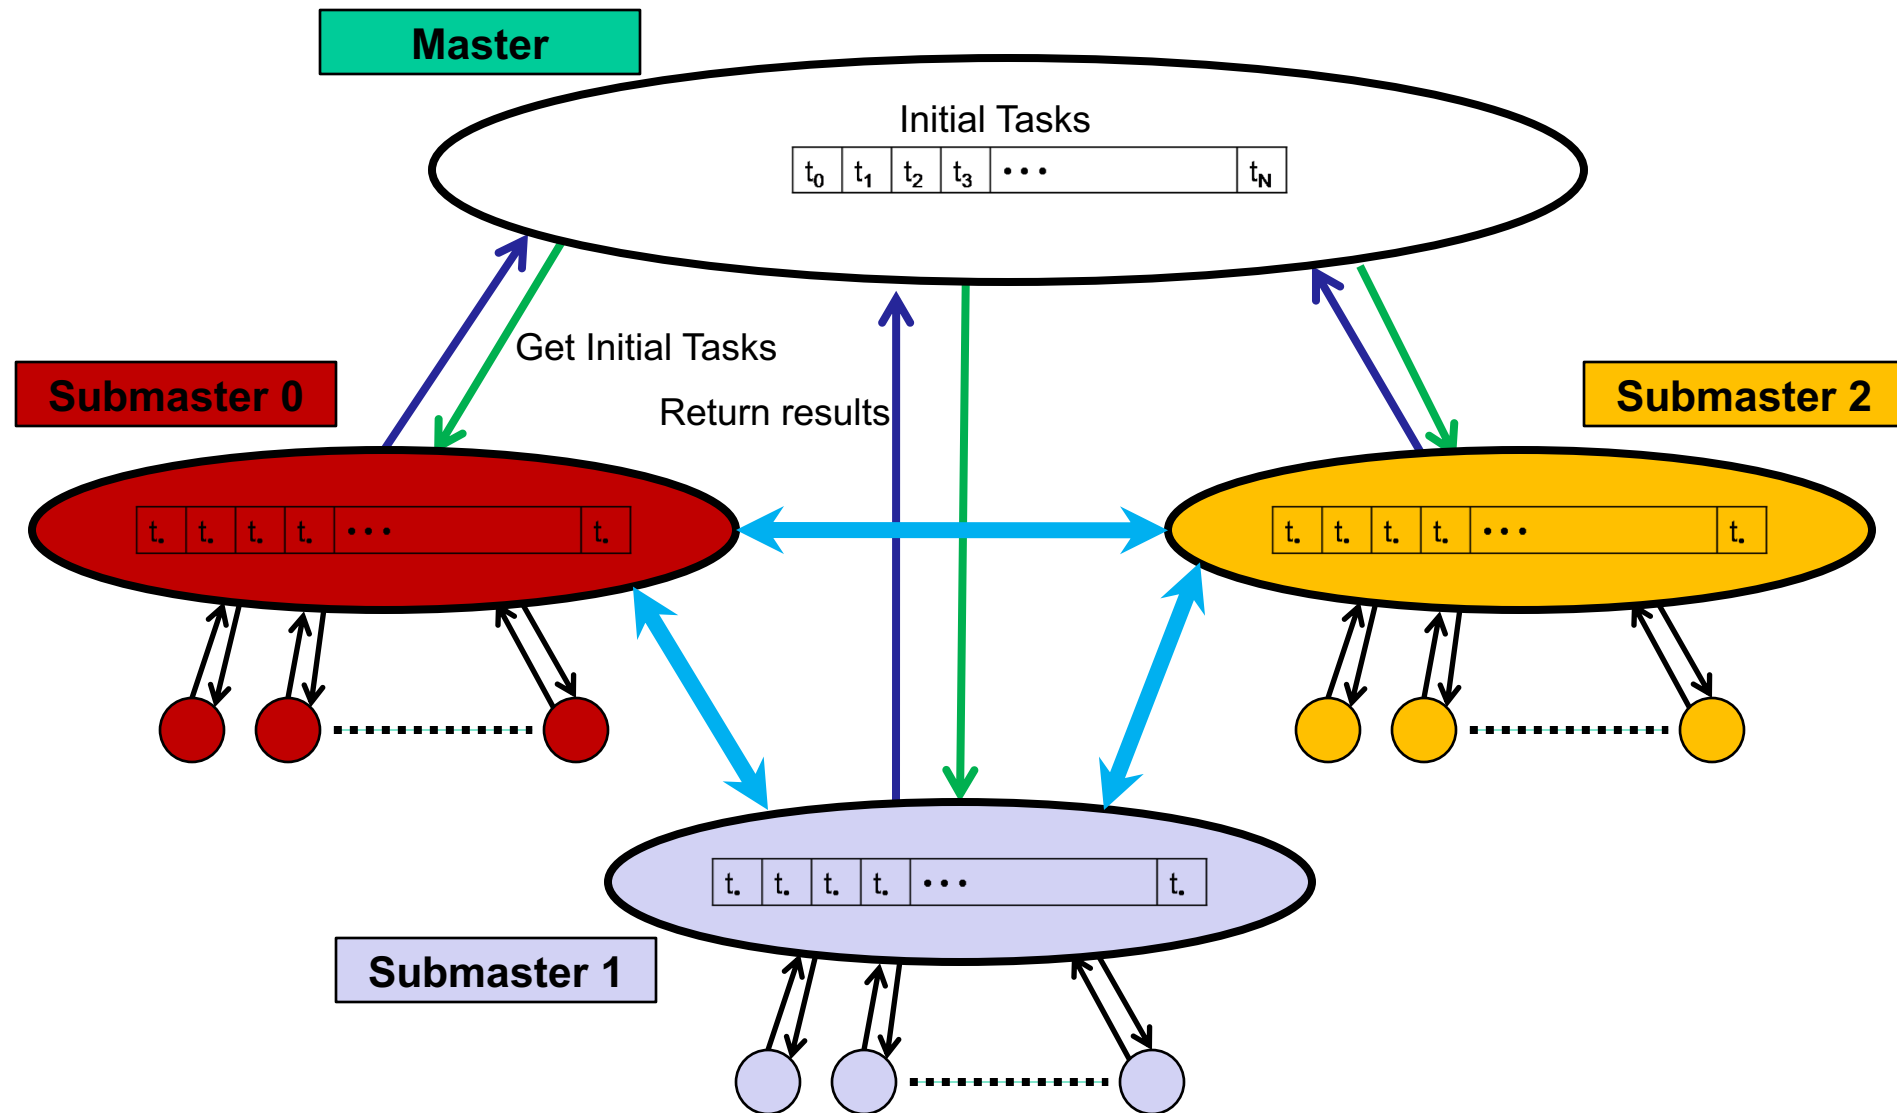


# Decentralized Task Pool

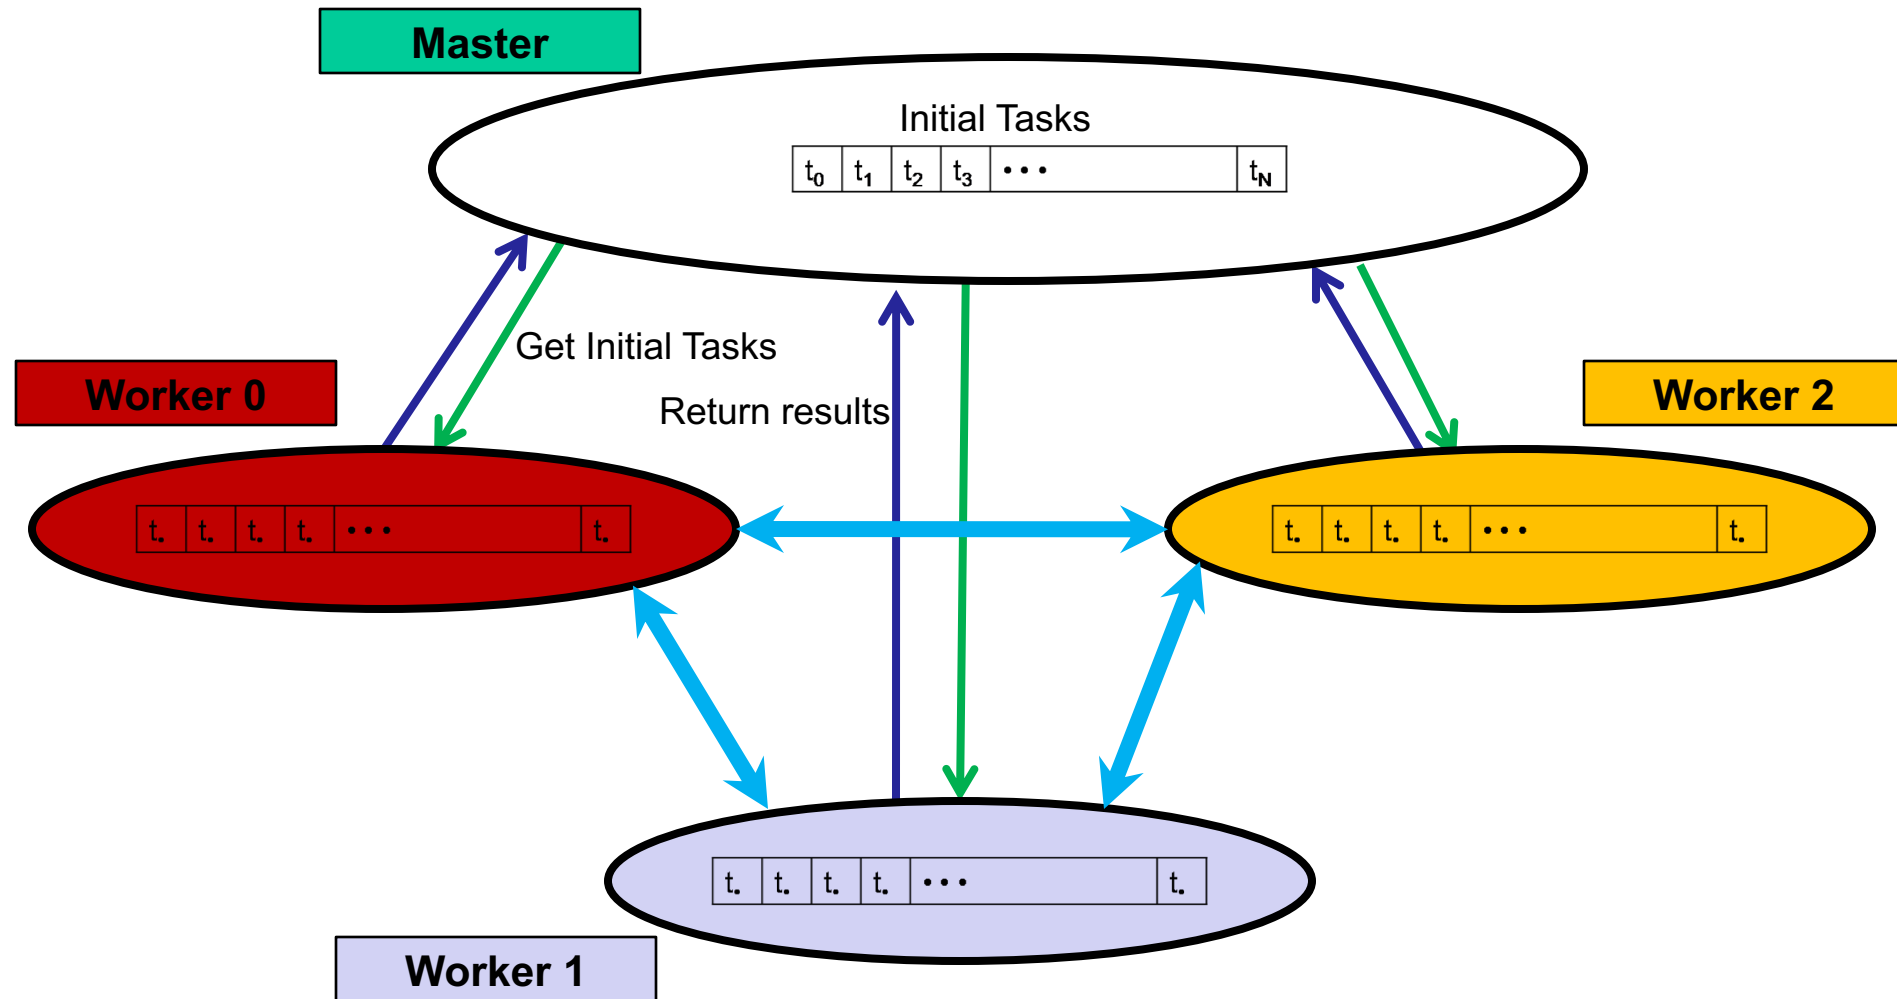
- The task pool is distributed among several processes, e.g.:
  - Hierarchy of “submasters,” each with separate worker sets, or
  - “Dual-mode” workers, acting as both submaster and worker



# Decentralized Task Pool Model (Submasters)



# Decentralized Task Pool Model (Dual-mode workers)



# Decentralized Task Pool


- The task pool is distributed among several processes, e.g.:
  - Hierarchy of “submasters,” each with separate worker sets, or
  - “Dual-mode” workers, acting as both submaster and worker
- Results all go back to the original master
- Termination: involves both local conditions (nothing to do right now) and global conditions (to guarantee that no additional tasks will show up)
- For the case of “dual-mode” workers:
  - Workers process tasks and possibly create new tasks
  - Workers pass tasks around among themselves to balance the load:
    - May “shed” (get rid of) tasks when they’re too busy
    - May request tasks when they’re free or don’t have many tasks queued up (like a watermarking approach)





# Task Transfer Mechanisms

Task pool is effectively “fully distributed” among the submasters or the dual-mode workers (which may include the master). The participants pass tasks among themselves to balance the load



# Task Transfer Mechanisms

- Receiver-Initiated
  - When its local task queue nears empty, worker requests tasks from one or more other workers that it selects
  - May work well when there are overloaded workers, but it is difficult to determine which workers those are
- Sender-Initiated
  - Workers “shed” work when they get backed up (lots of tasks waiting in their local queues)
  - Generally works well if some workers have light loads, but, again, it may be difficult to determine which ones those are.
- Practical implementations may combine the above and may use a “load-independent task” movement algorithm
- Some systems use a “bulletin board” model—possibly hosted by the master (e.g. “Linda” or database systems, where all tasks are visible via query)




# Worker Selection for Task Transfer

- Round Robin
  - Each worker transfers tasks to or from a process it selects using a counter that rotates around all the workers
    - Initially: *counter* for  $p_i$  is  $p_{i+1} \pmod{nw}$  for  $nw$  workers
    - After an attempt to transfer:  $counter = counter + 1 \pmod{nw}$ , leaving out the process itself
- Random Polling
  - Worker  $p_i$  selects worker  $p_x$  where  $x$  is an integer in  $[0, nw-1]$ , excluding  $i$
- Fixed Topology Task Shifting
  - Workers receive tasks from specific worker(s) and pass them on until they reach lightly-loaded worker(s)




# Line Structure Task Shifting

Master feeds tasks to  $p_1$ , and the tasks are then shifted down the line to idle workers.



Possible implementation might use two threads or two processes per worker (on a single cpu): one for communication about tasks, and one for real work



# Distributed Termination Detection

- Conditions:
  - Application-specific local termination conditions on each worker
  - No task assignments pending or “in transit” (difficult to know)
- More general and safer approach:
  - Each worker has 2 states: “*Inactive*” (ready to terminate) or “*Active*.” Worker starts *Inactive* and becomes *Active* when it first receives a task. Source of that task becomes the worker’s “*Parent*”
  - Worker immediately acknowledges all tasks received, except its initial task, which is only acknowledged just before it becomes inactive
  - Once active, a worker stays active until it:
    - meets local termination conditions and has no pending tasks
    - has received all acknowledgements due for tasks it sent out
    - has acknowledged all tasks it received other than the first one
    - acknowledges initial task to its parent (last step before becoming inactive)
  - Termination occurs when master becomes inactive



# Termination Using Acknowledgements

