



Partitioning and Divide-and-Conquer Strategies – Part I

CPSC 424/524
Lecture #10
October 31, 2018

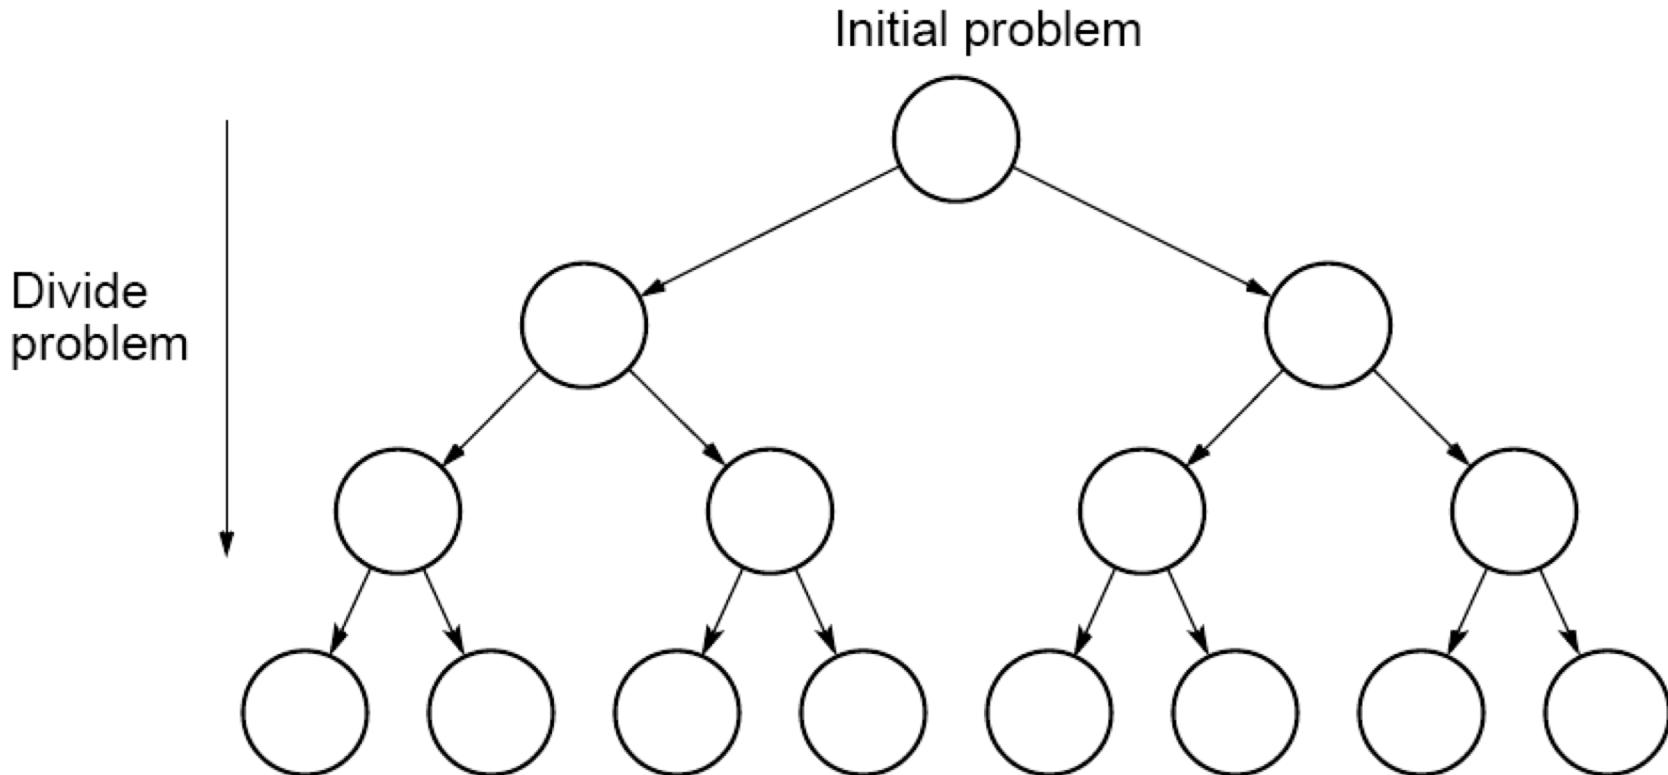


Partitioning vs. Divide-and-Conquer

- Partitioning
 - Divides a problem into parts
 - Solves the individual parts
 - Combines the results
 - Parts may be completely dissimilar
- (Recursive) Divide-and-Conquer
 - Special case of partitioning
 - Partitioned sub-problems have same form as the original problem
 - Partitioning repeated recursively; often represented in tree form
 - Ideally, data usage for each sub-problem is localized



Recursive Partitioning



Examples

- Arithmetic operations on sequences of numbers (Reductions)
- Matrix multiplication in block form
- Sorting numbers (more later)
- Numerical Integration (more later)
- N-body problem (more later)
- Database searching
- Sequence comparison
- Grid computations (e.g. Partial differential equations)



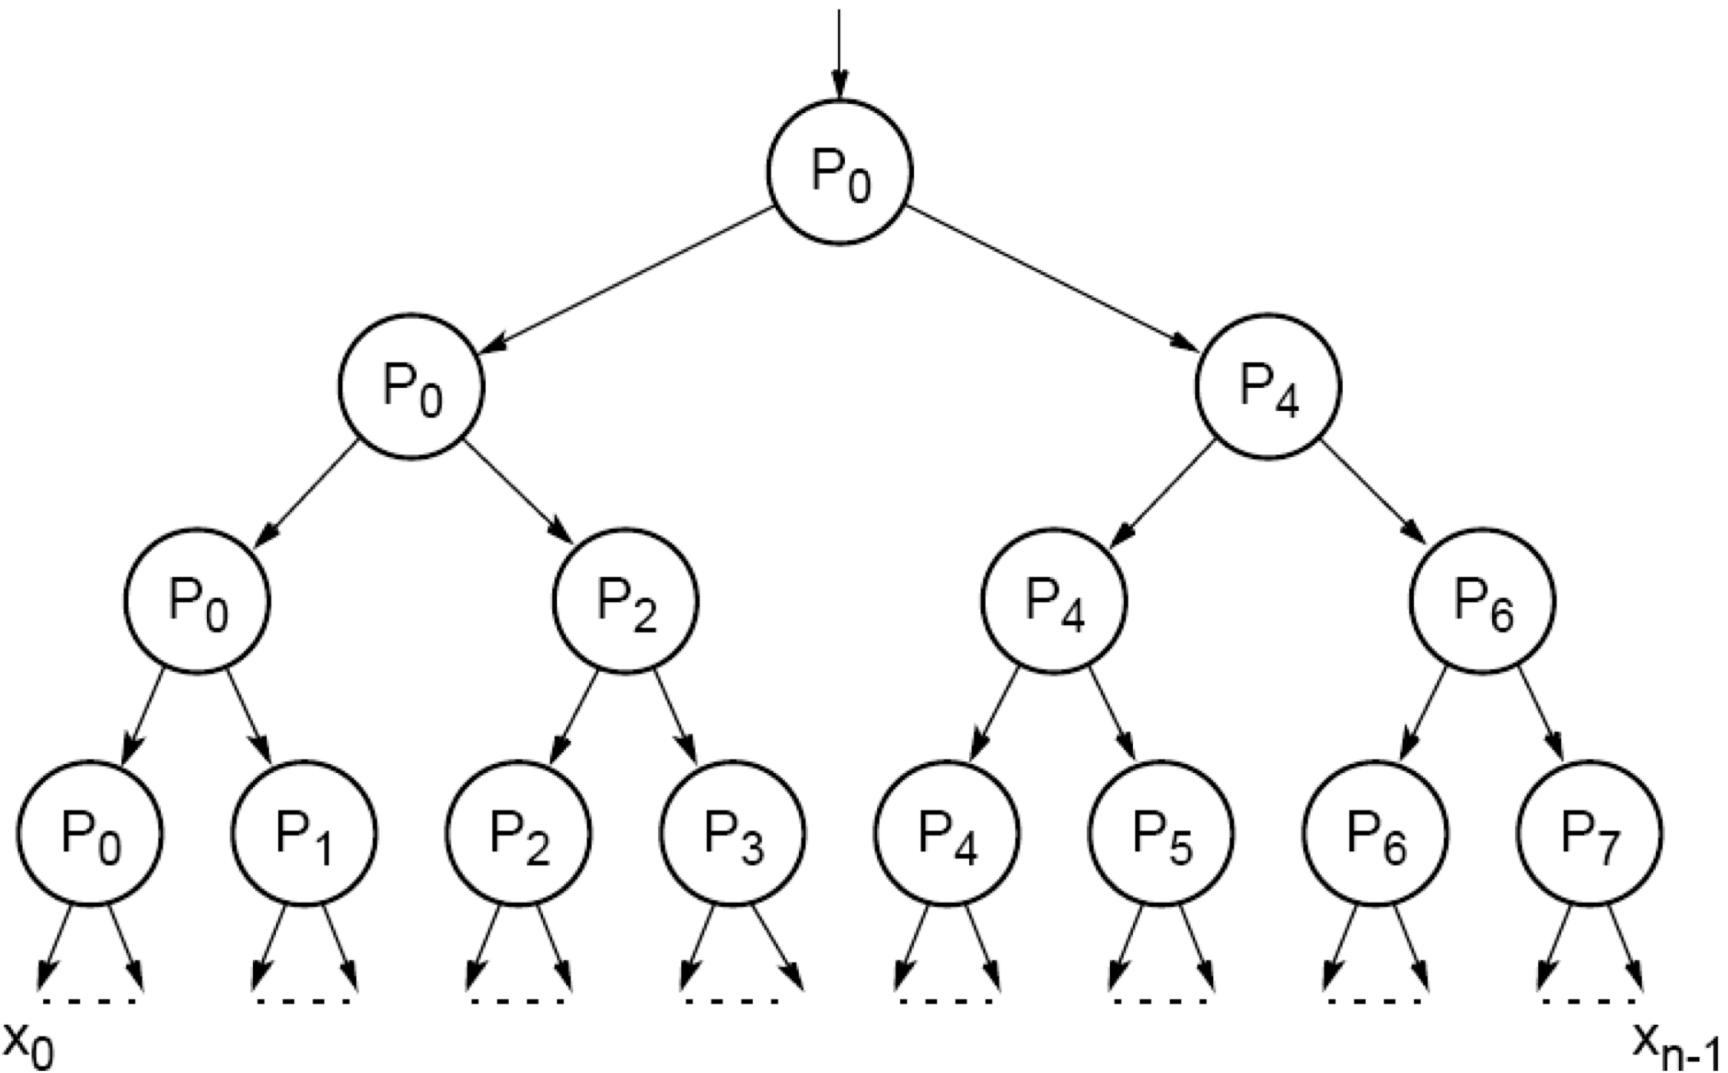
Example: Summing a sequence of numbers

$$S = \sum_{i=0}^{n-1} x_i$$

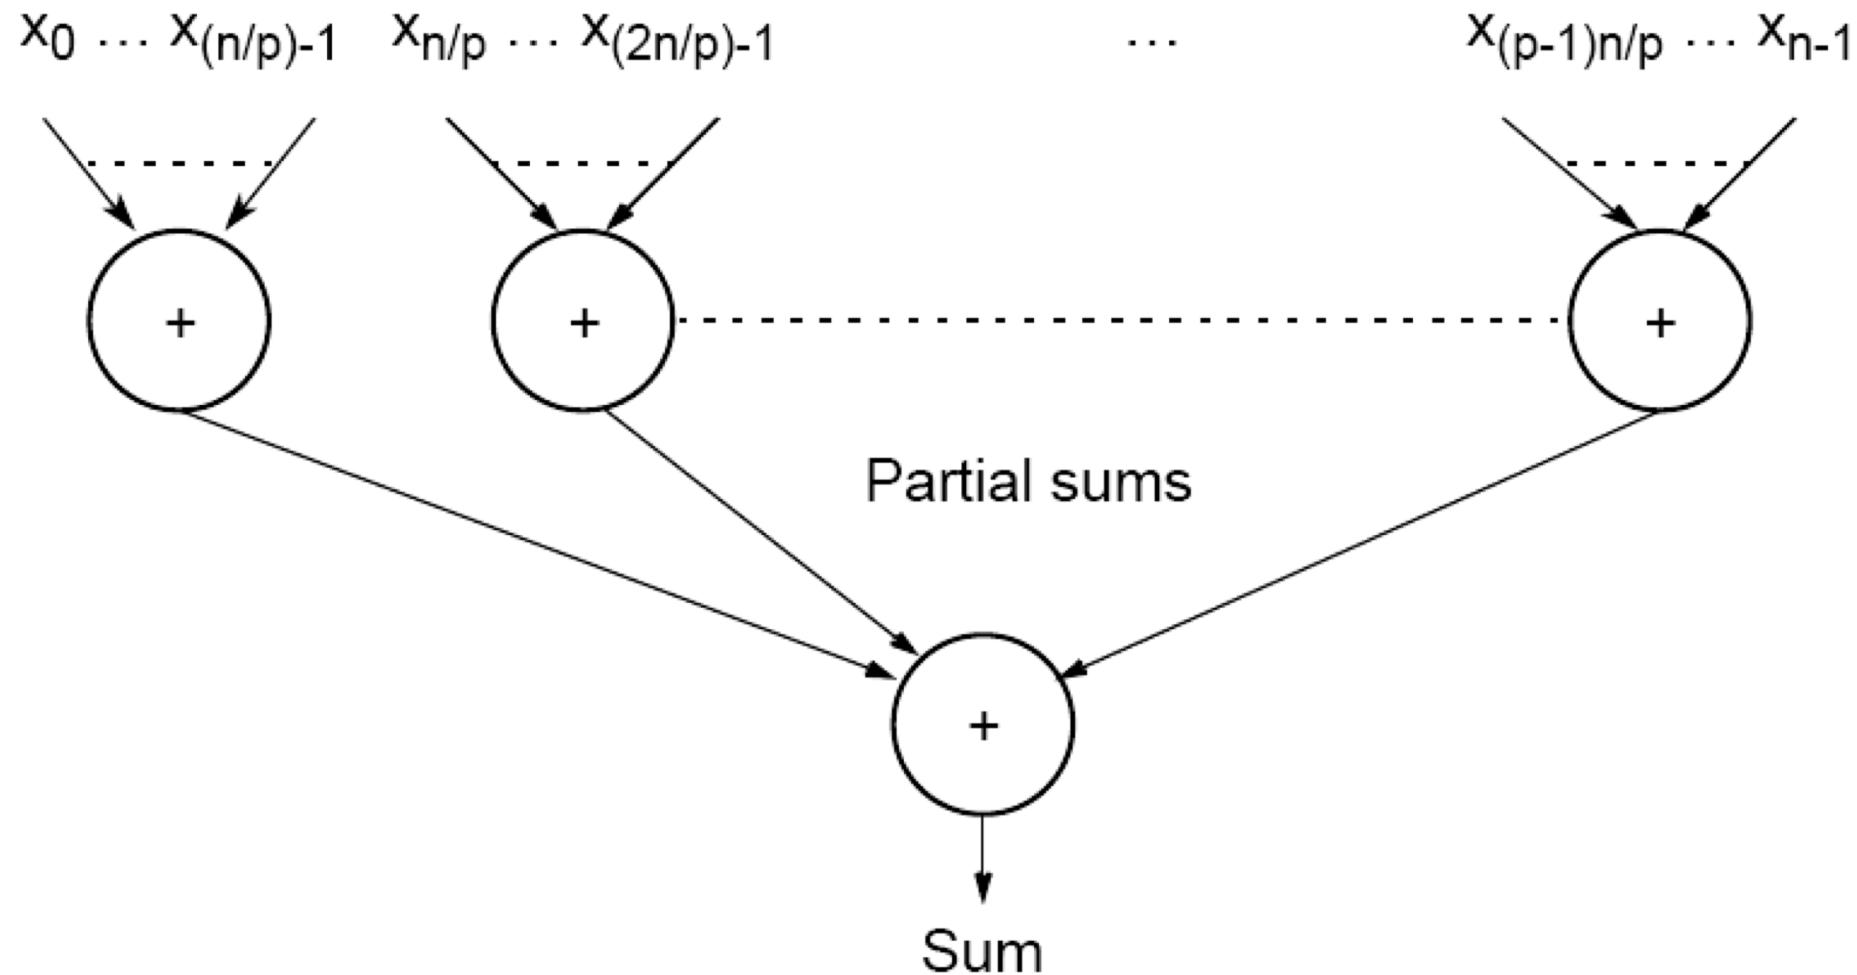


Partitioning the Sequence

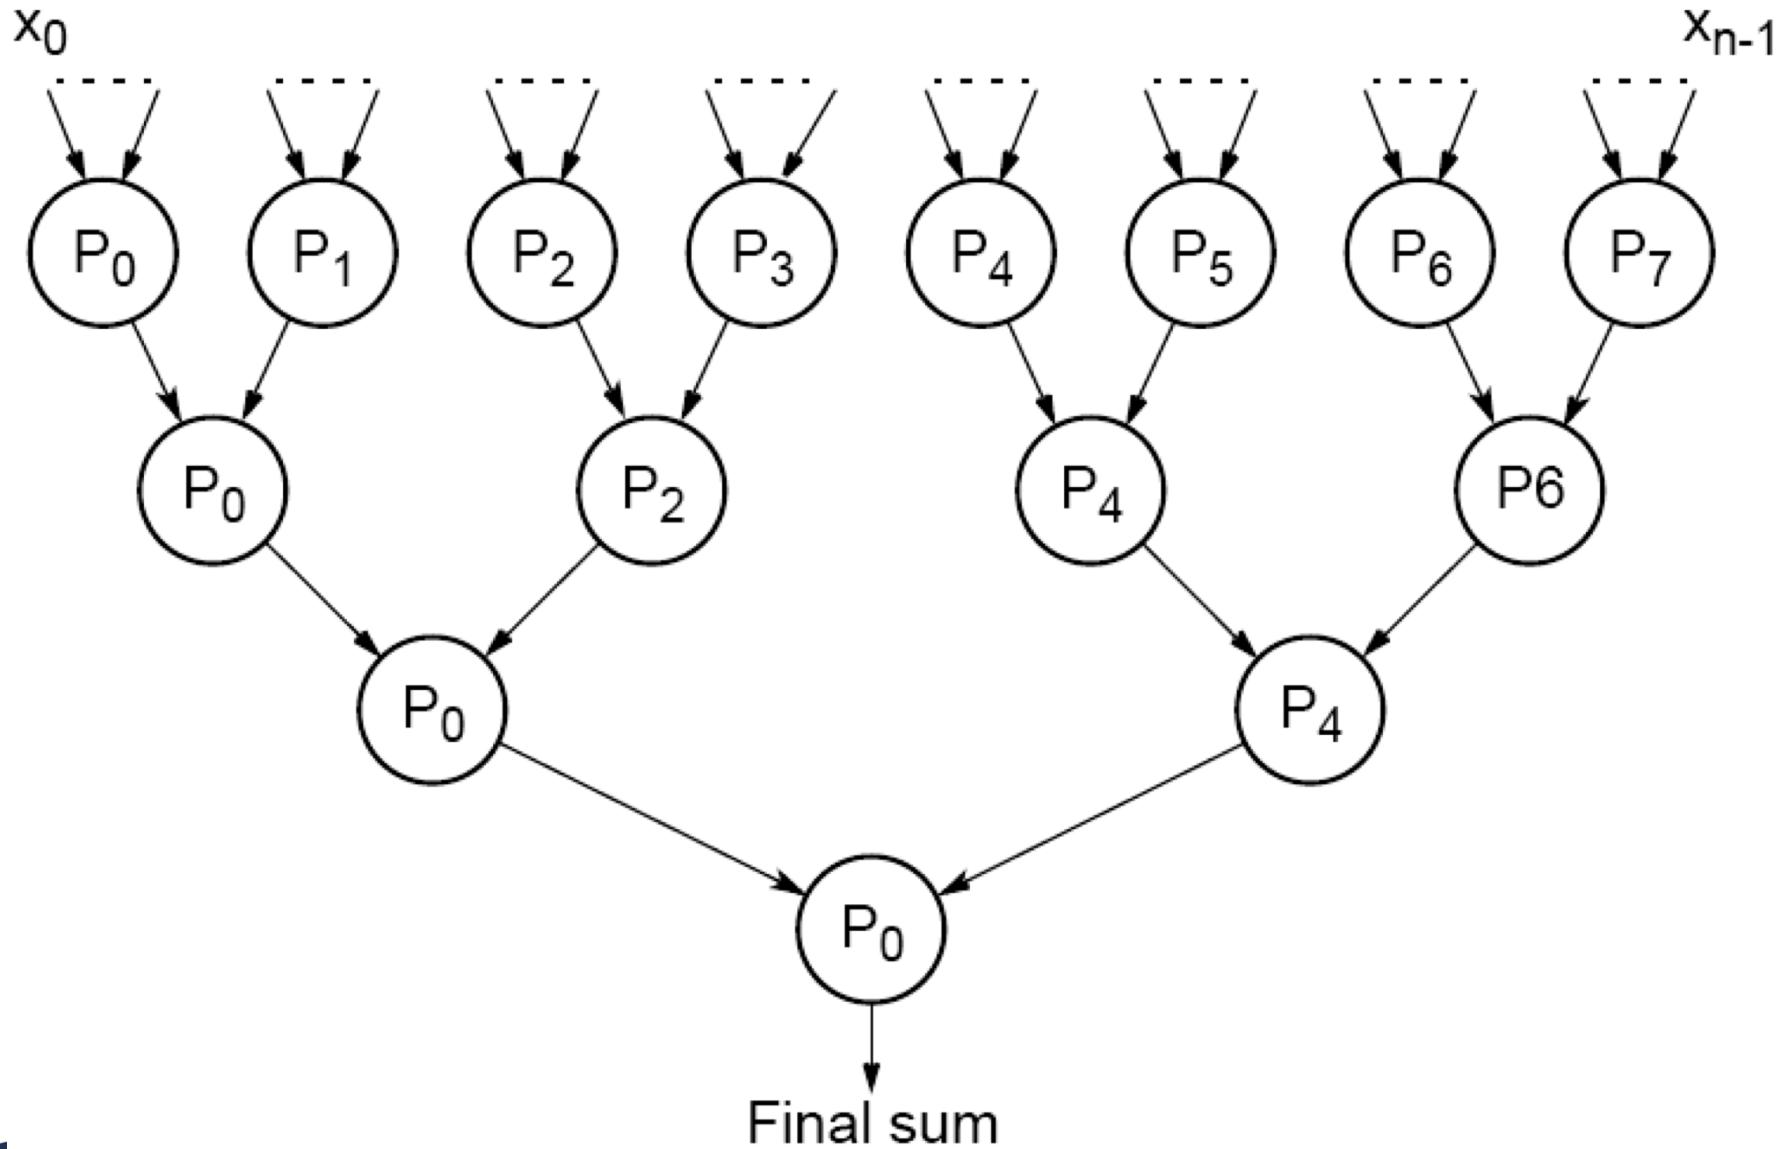
Original Sequence: $x_0, x_1, x_2, \dots, x_{n-1}$



Summing the Partitioned Sequence



Carrying Out the Summation



Sorting Algorithms

Many sorting algorithms can be parallelized using divide-and-conquer.

Examples:

- Bucket sort
- Quicksort

Start with n numbers (assume no duplicates)

Select one at random, say x_s

Split the problem into two parts:

$S_<$: those numbers $< x_s$

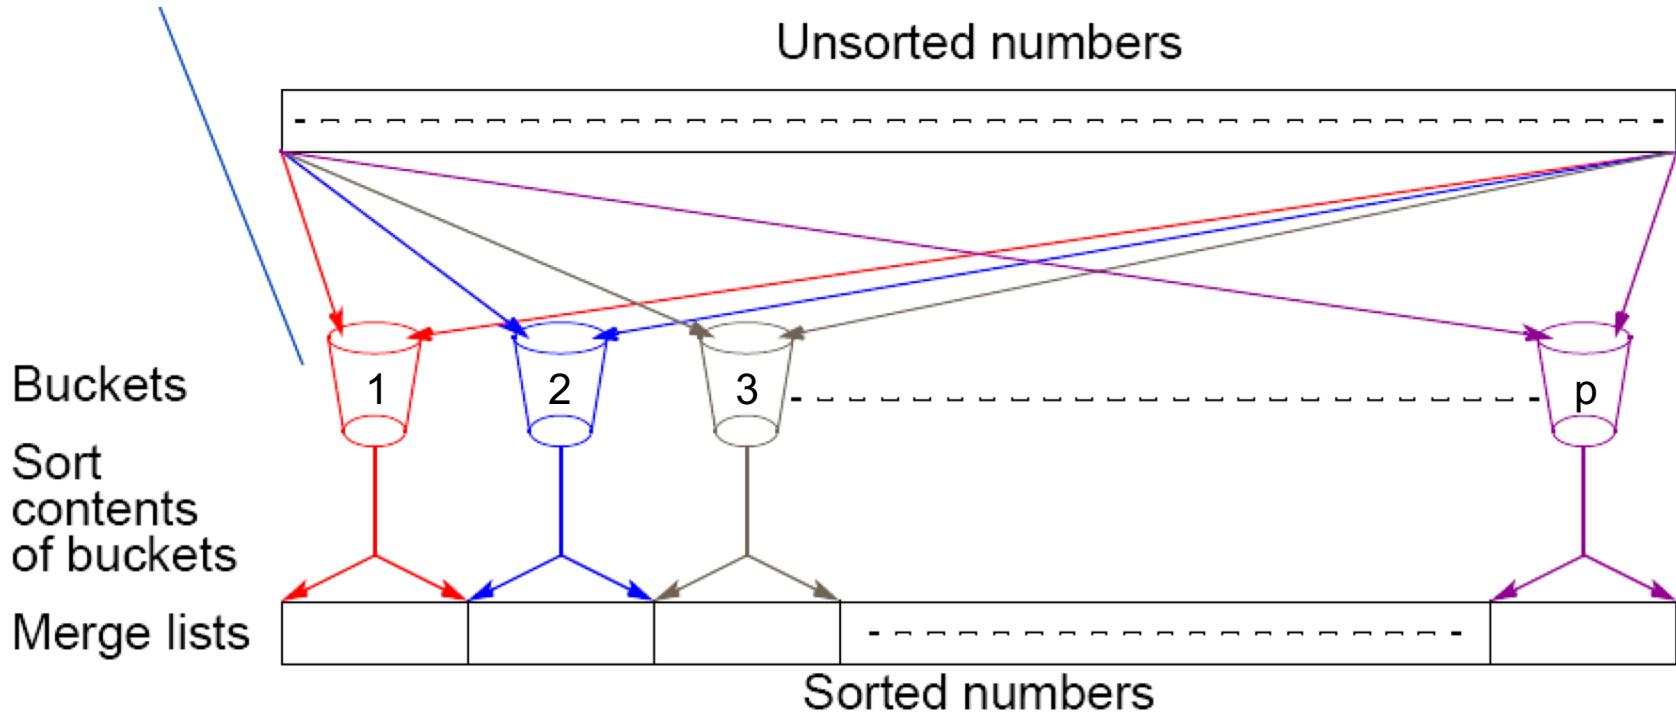
$S_>$: those numbers $> x_s$

Repeat until the subproblems are small enough



Bucket Sort

One “bucket” assigned to hold numbers that fall within each region.
Numbers in each bucket sorted using a sequential sorting algorithm.

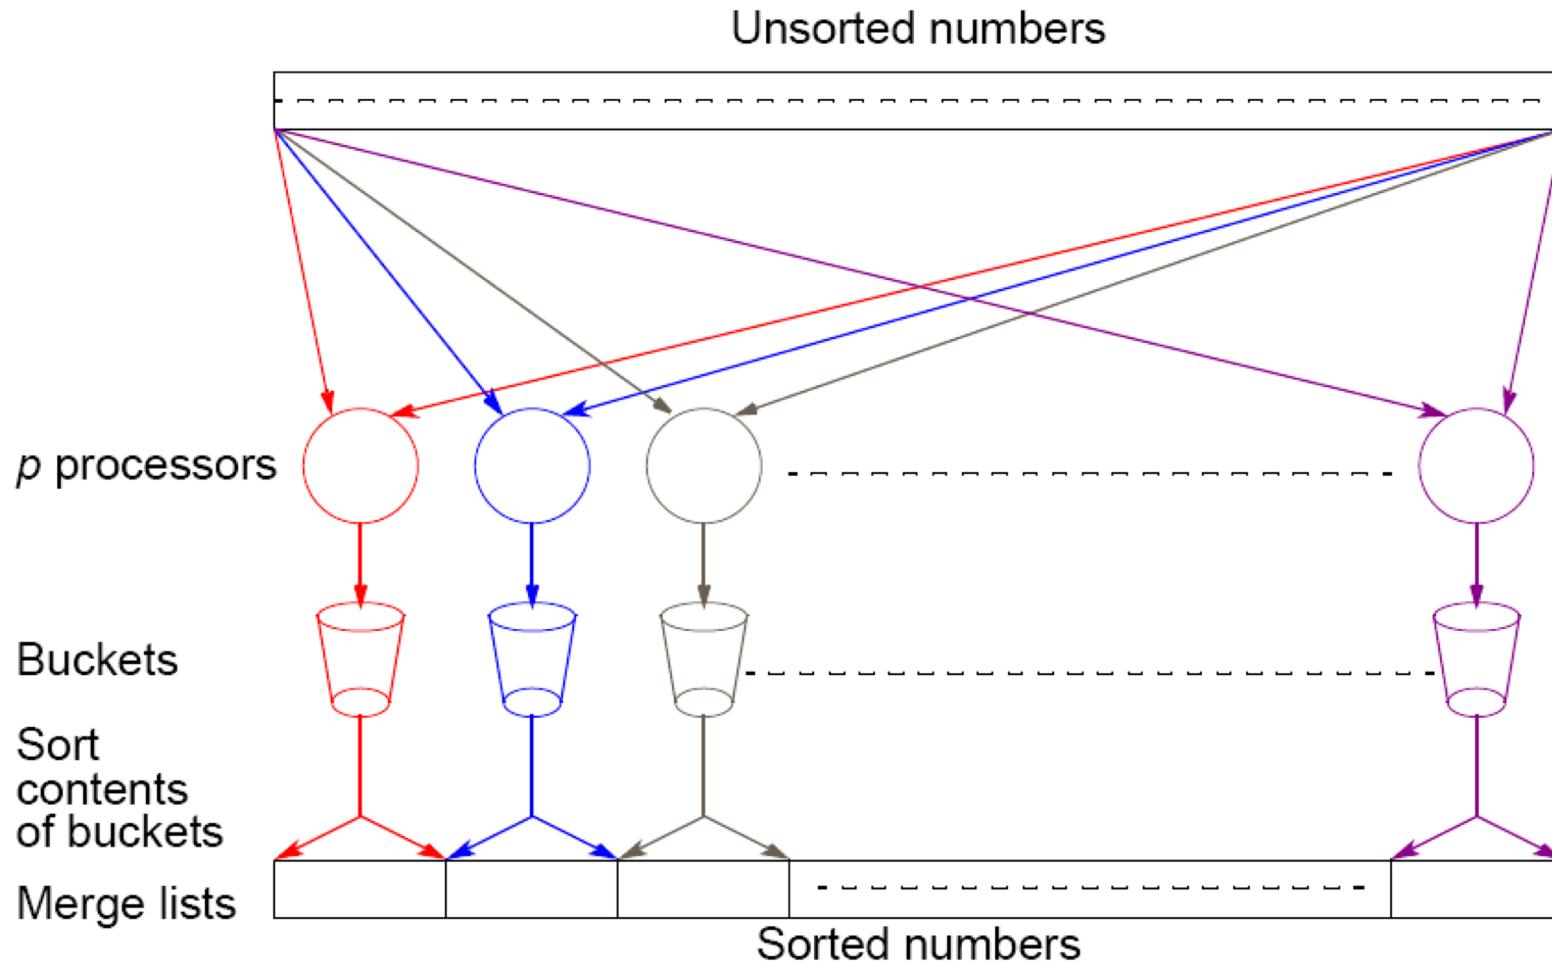


Sequential sorting time: $O(p \cdot (n/p) \cdot \log(n/p))$ for p buckets (best case).
Works well if the original numbers uniformly distributed across a known interval, say 0 to $M-1$.



Simple Parallel Version of Bucket Sort

Simple Approach: Assign one processor for each bucket.
(Each process looks at all numbers in first phase.)



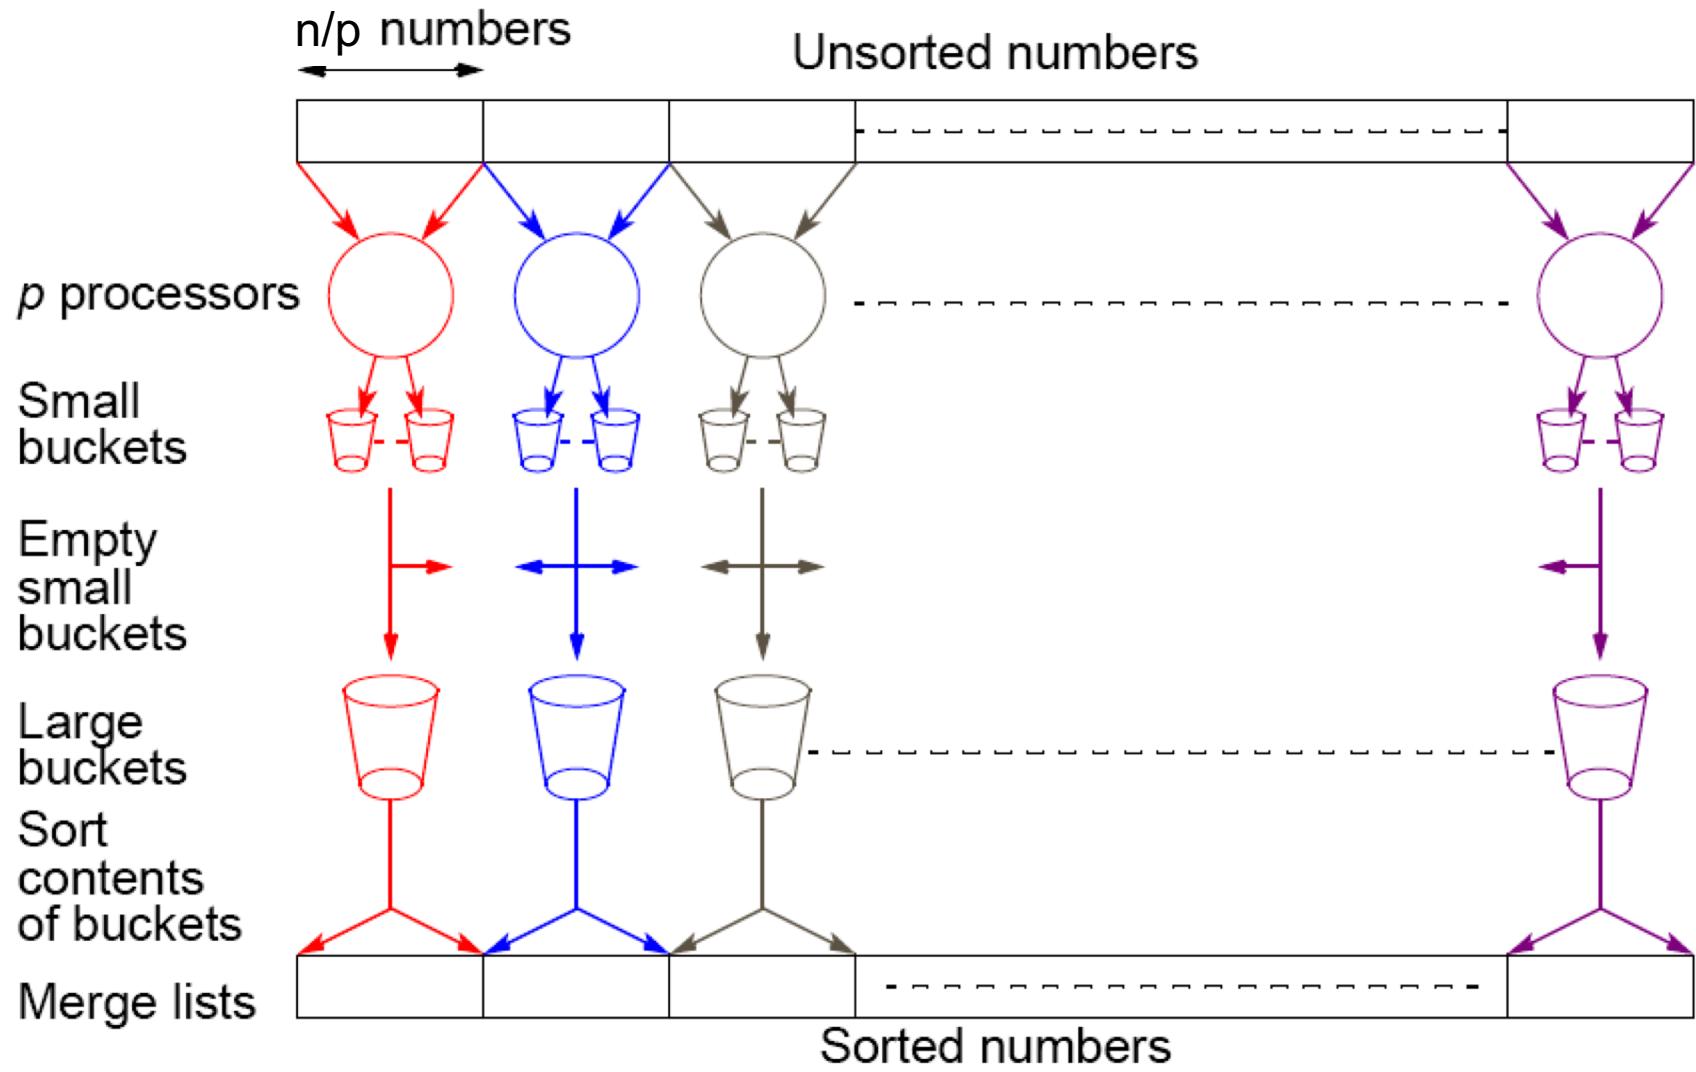
Better Parallelization

Using p processors to sort n numbers:

- Partition sequence range into p disjoint intervals, 1 per processor
- Give each processor $\sim n/p$ (randomly selected) numbers
- Each processor “deals” its numbers into local buckets:
 - A “large” bucket for its own interval (which it keeps)
 - $p-1$ “small” buckets for other intervals
(which will be sent to the processors responsible for those intervals)
- Processors exchange the small buckets
- Each processor sorts all the numbers it has
- Master merges results
 - Easy, since sorted buckets correspond to disjoint, ordered intervals



Parallel Bucket Sort

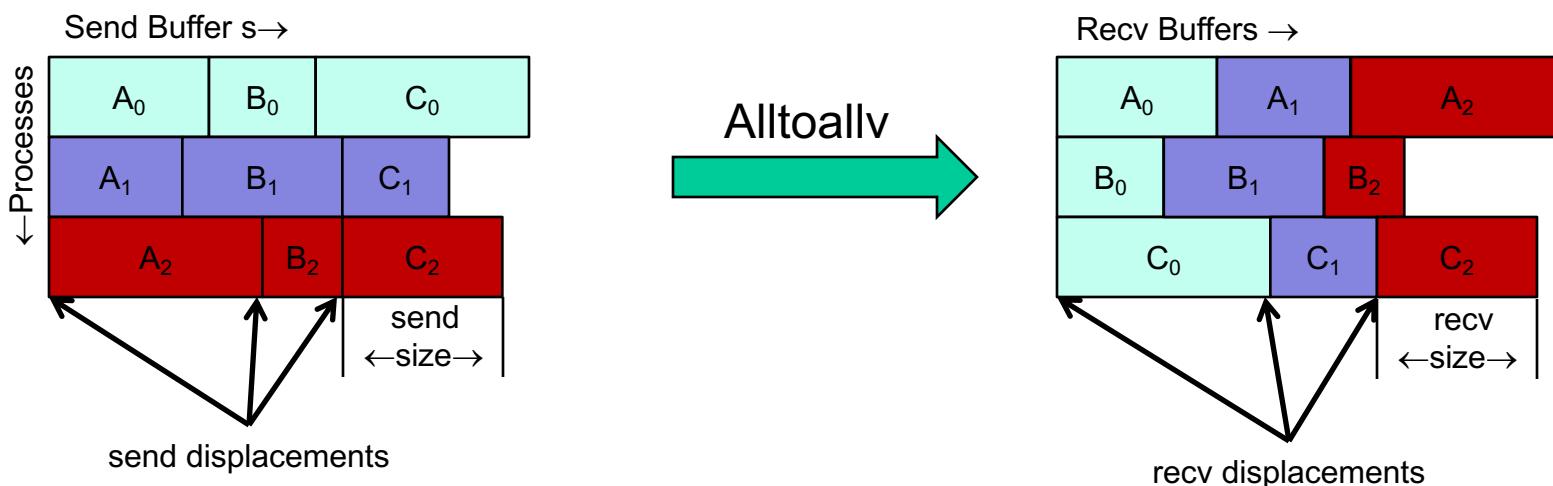
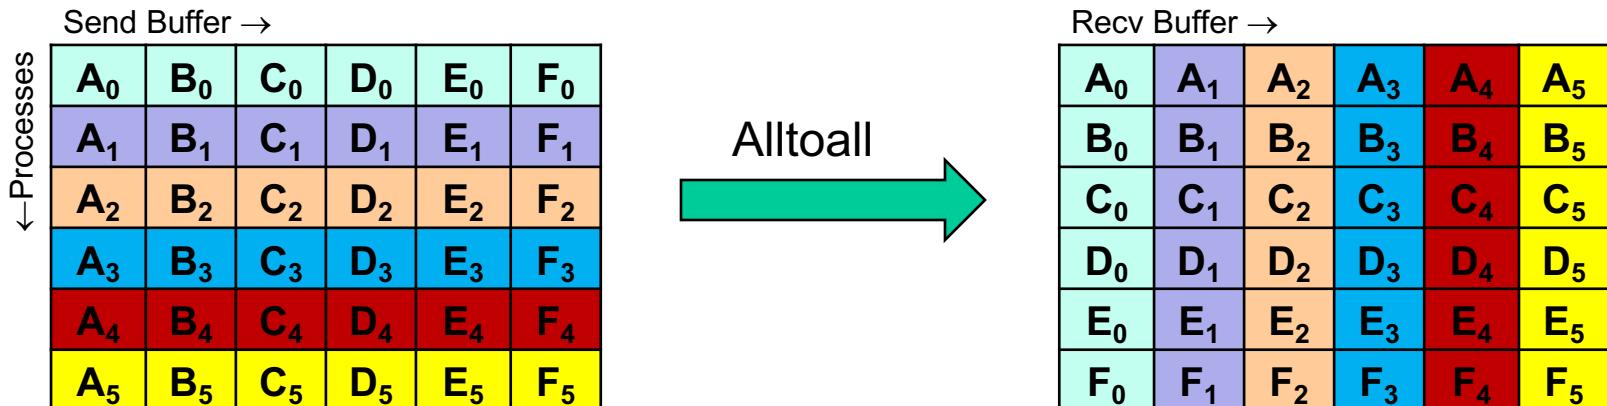


Generalized Gather/Scatter-like Operations

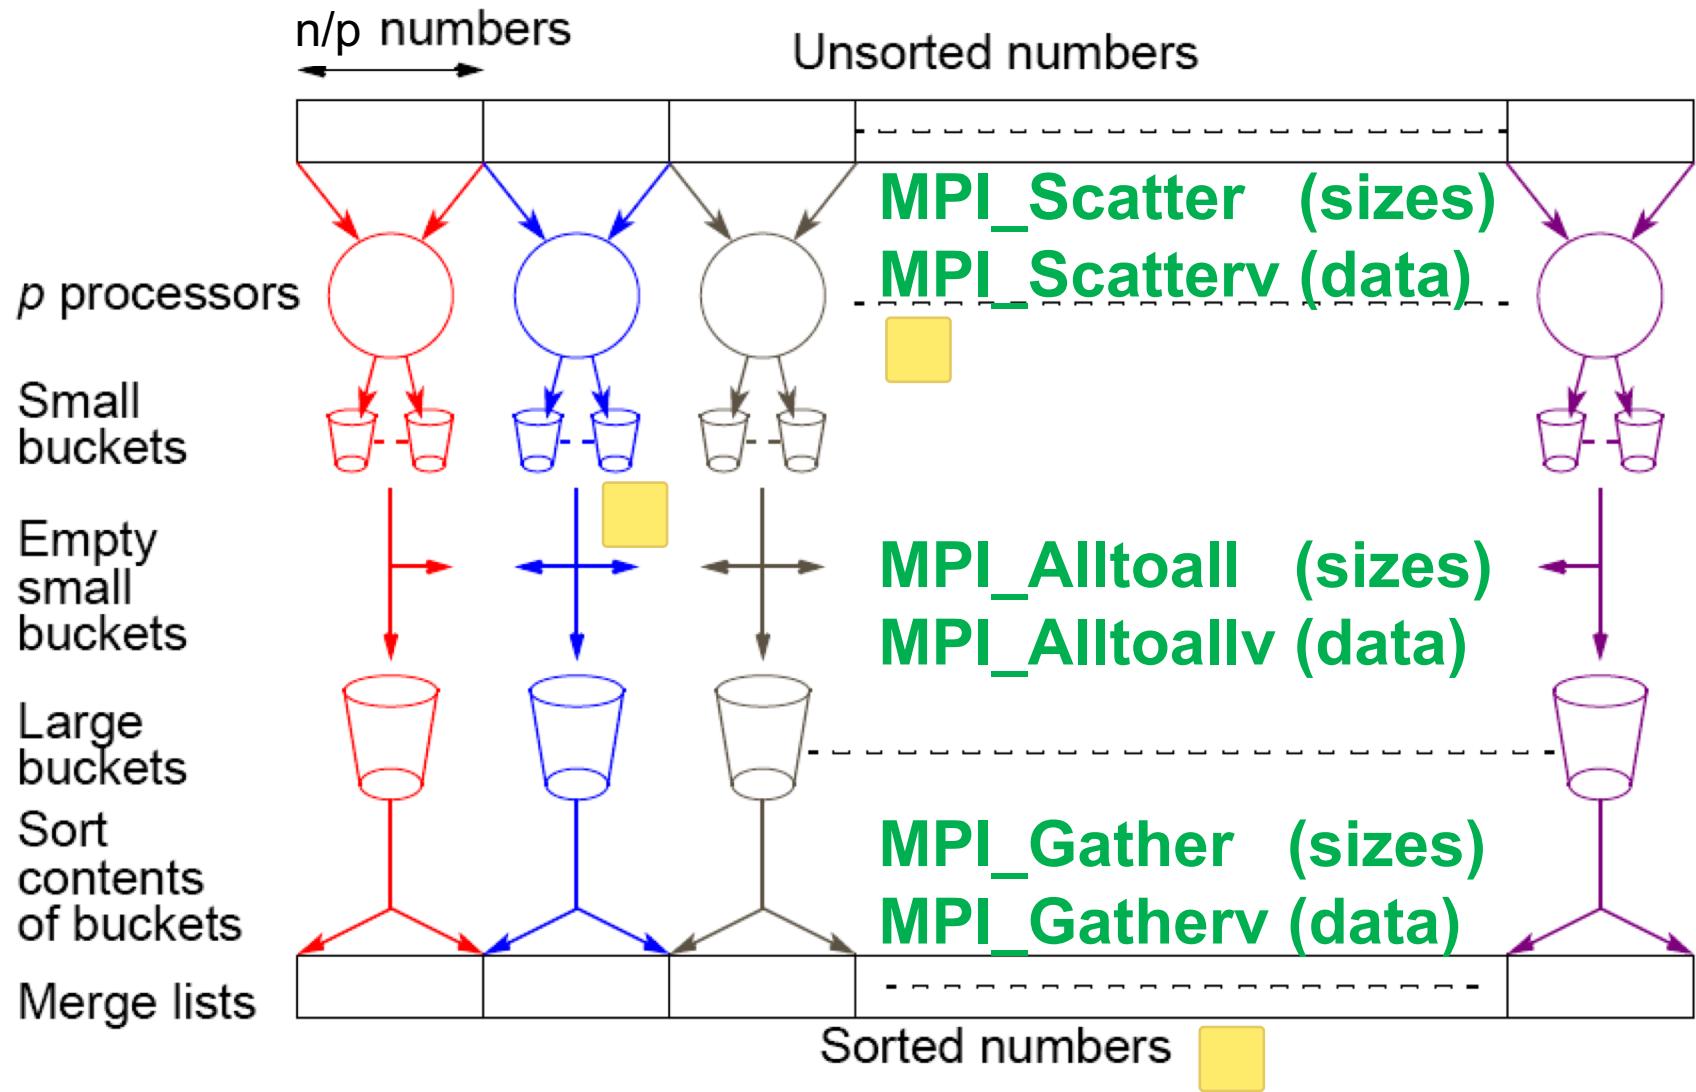
- Vector versions: Permit varying segment sizes & displacements
 - `MPI_Scatterv`
 - `MPI_Gatherv`
 - `MPI_Allgatherv`
 - `MPI_Alltoallv`
- “W” Version: Extends “V” version with varying data types
 - `MPI_Alltoallw`
 - Can be specialized to act like “W” versions of other collectives
- Intent is to allow flexibility in the location of the data on the send and/or recv side of the operations. (Obviously useful in cases like bucket sort.)



Alltoall and Alltoally



Alternative Bucket Sort



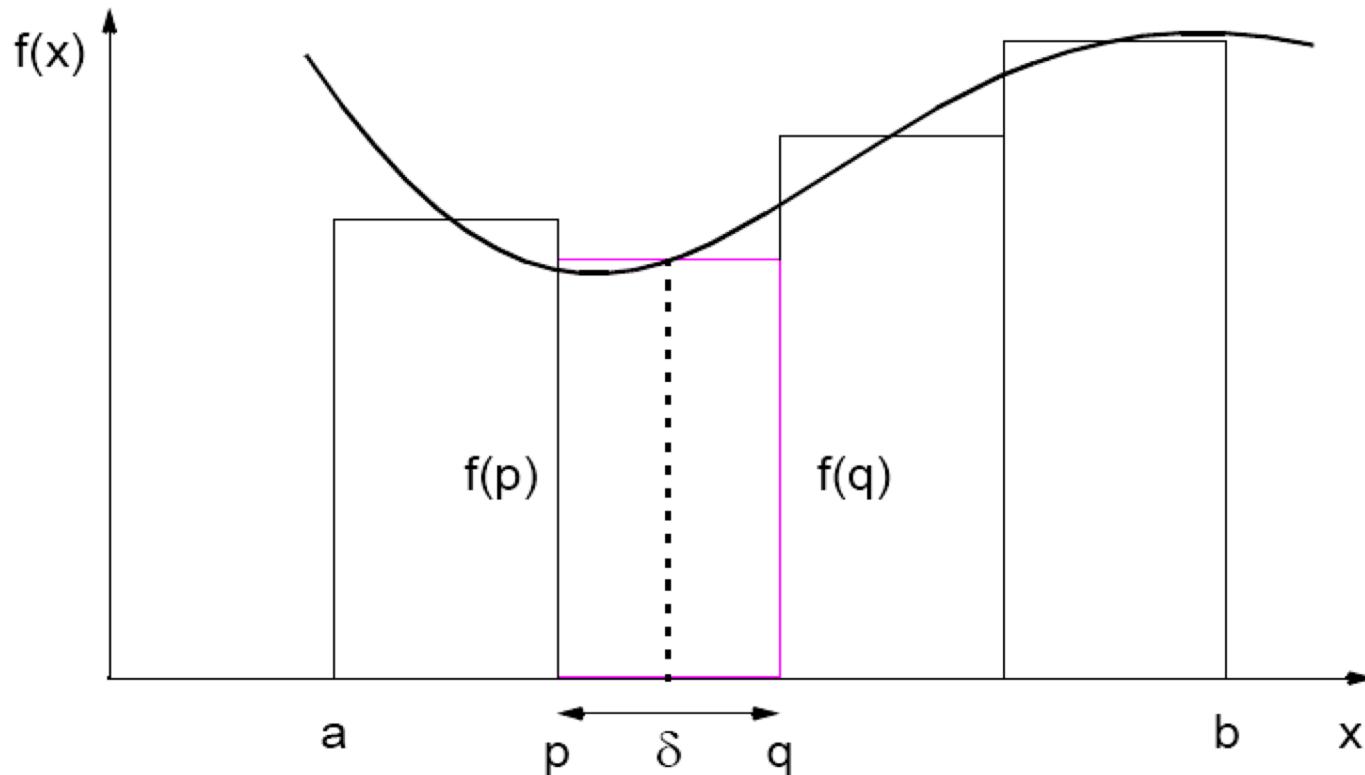
Collectives for Bucket Sort

```
int size, *recvbuf, *sendbuf;  
int sendsizes[p], senddisps[p], recvsizes[p], recvdisps[p];  
.  
.  
.  
//sendsizes[i] = count of data items to be sent to process i  
//  
//  
//senddisps[i] = displacement in sendbuf of data for process i  
//  
//  
//computed based on sendsizes)  
  
MPI_Alltoall(sendsizes, 1, MPI_INT, recvsizes, 1, MPI_INT,  
              MPI_COMM_WORLD);  
  
//recvdisps[i] is displacement in recvbuf for data from proc i  
for (i=1,recvdisps[0]=0; i<p; i++)  
    recvdisps[i] = recvdisps[i-1] + recvsizes[i-1];  
  
MPI_Alltoallv(sendbuf, sendsizes, senddisps, MPI_INT,  
              recvbuf, recvsizes, recvdisps, MPI_INT,  
              MPI_COMM_WORLD);
```

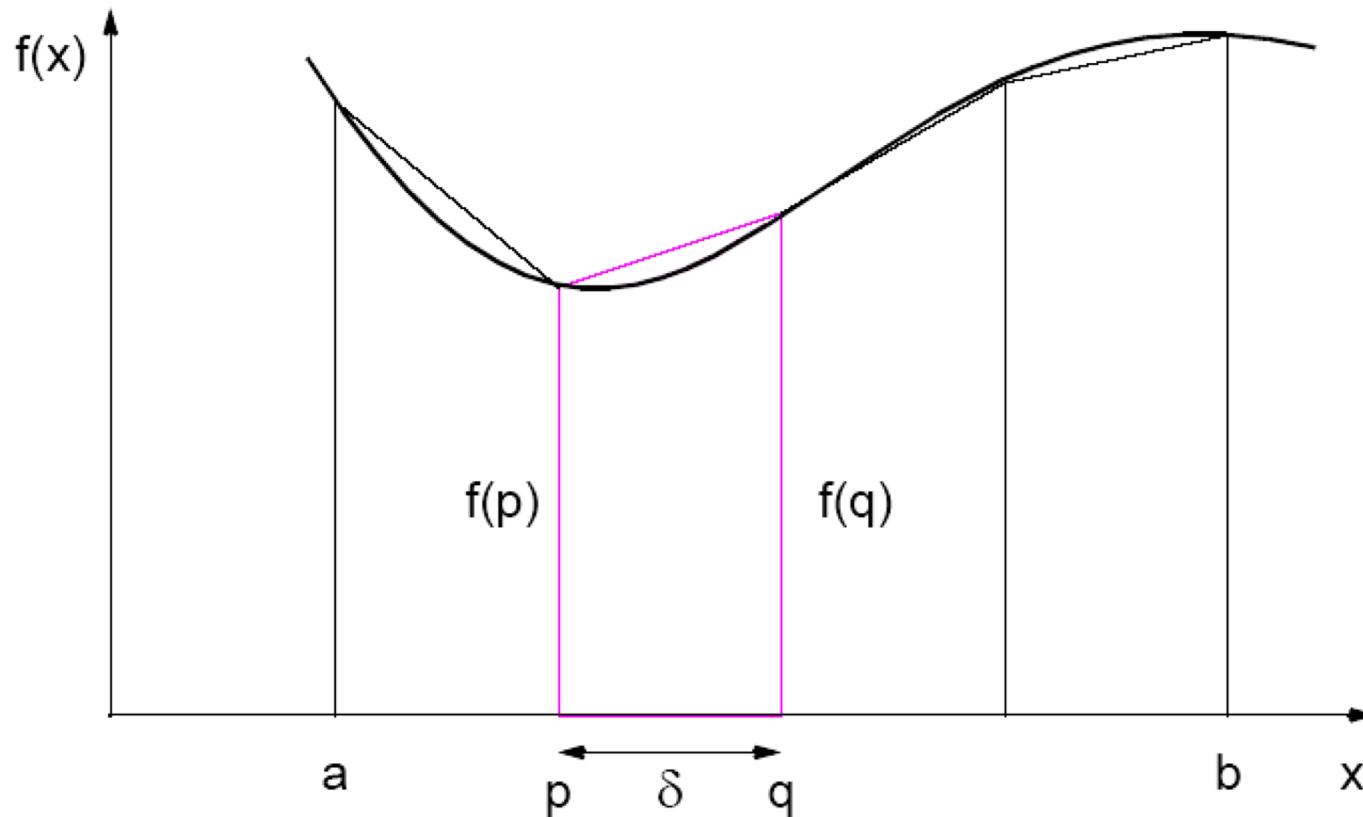


Numerical integration using midpoint rule

Each region calculated using an approximation given by rectangles:

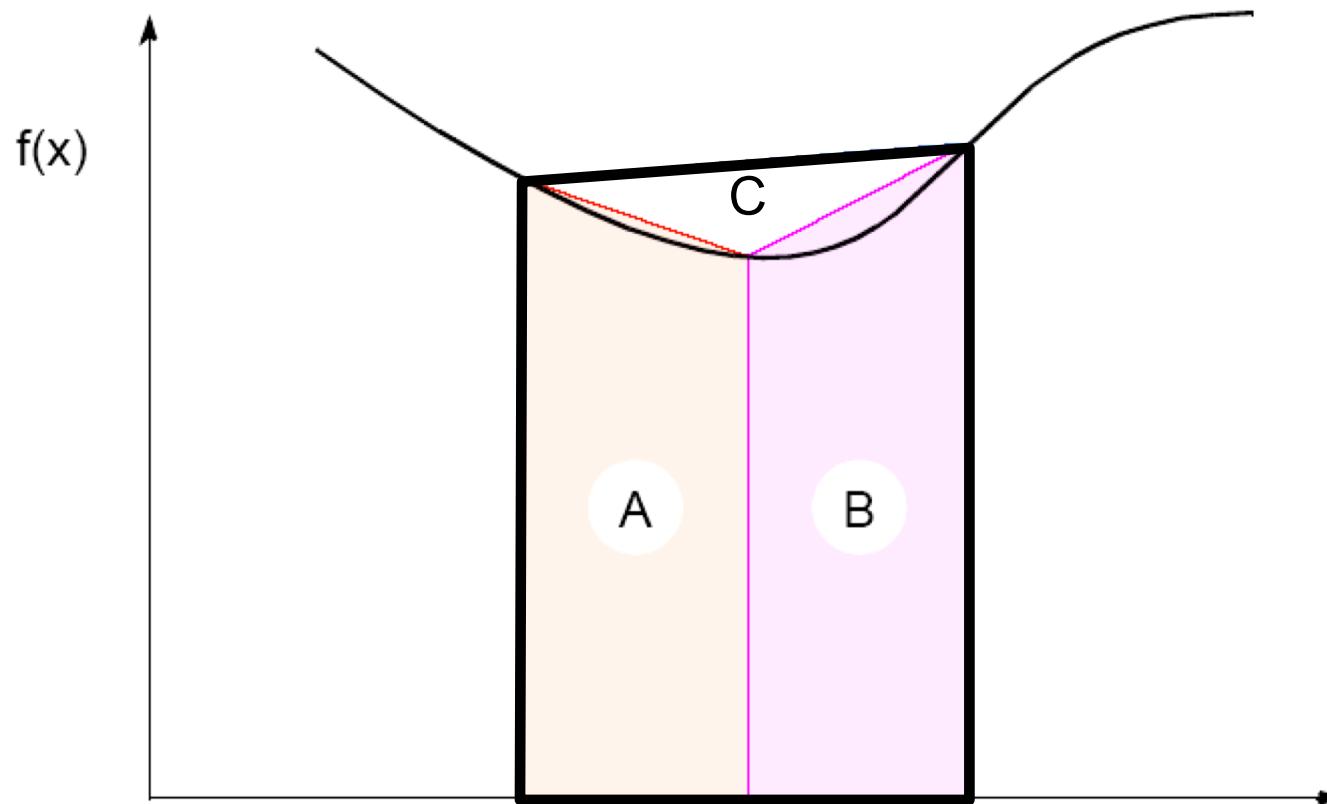


Numerical integration using trapezoidal rule



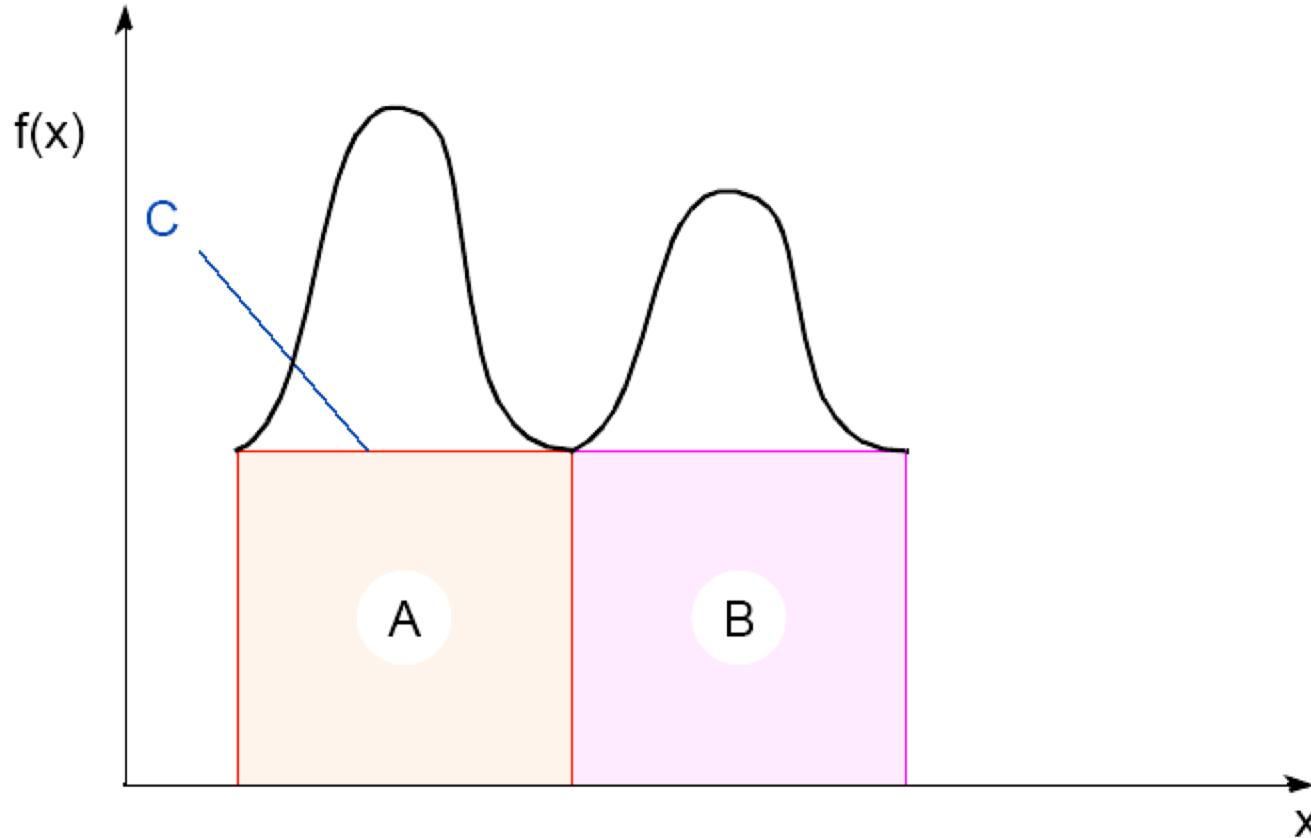
Adaptive Quadrature

Solution adapts to shape of curve. Use three areas, A , B , and C (from the previous interval). Computation terminated when $A + B$ sufficiently close to C .



Termination Issues

Some care might be needed in choosing when to terminate.



Might cause early termination, as two large regions are the same (i.e., $C = A+B$).

